

 Code

```
!unzip ibm-hr-analytics-attribution-dataset.zip
```



Archive: ibm-hr-analytics-attribution-dataset.zip
inflating: WA_Fn-UseC_-HR-Employee-Attrition.csv

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import pylab as pl
from sklearn.metrics import roc_curve, auc
from sklearn import metrics
from sklearn import feature_selection
import itertools
from sklearn.model_selection import cross_val_score
import warnings
import seaborn as sns
warnings.filterwarnings("ignore")

df = pd.read_csv('WA_Fn-UseC_-HR-Employee-Attrition.csv')
df.head()
```




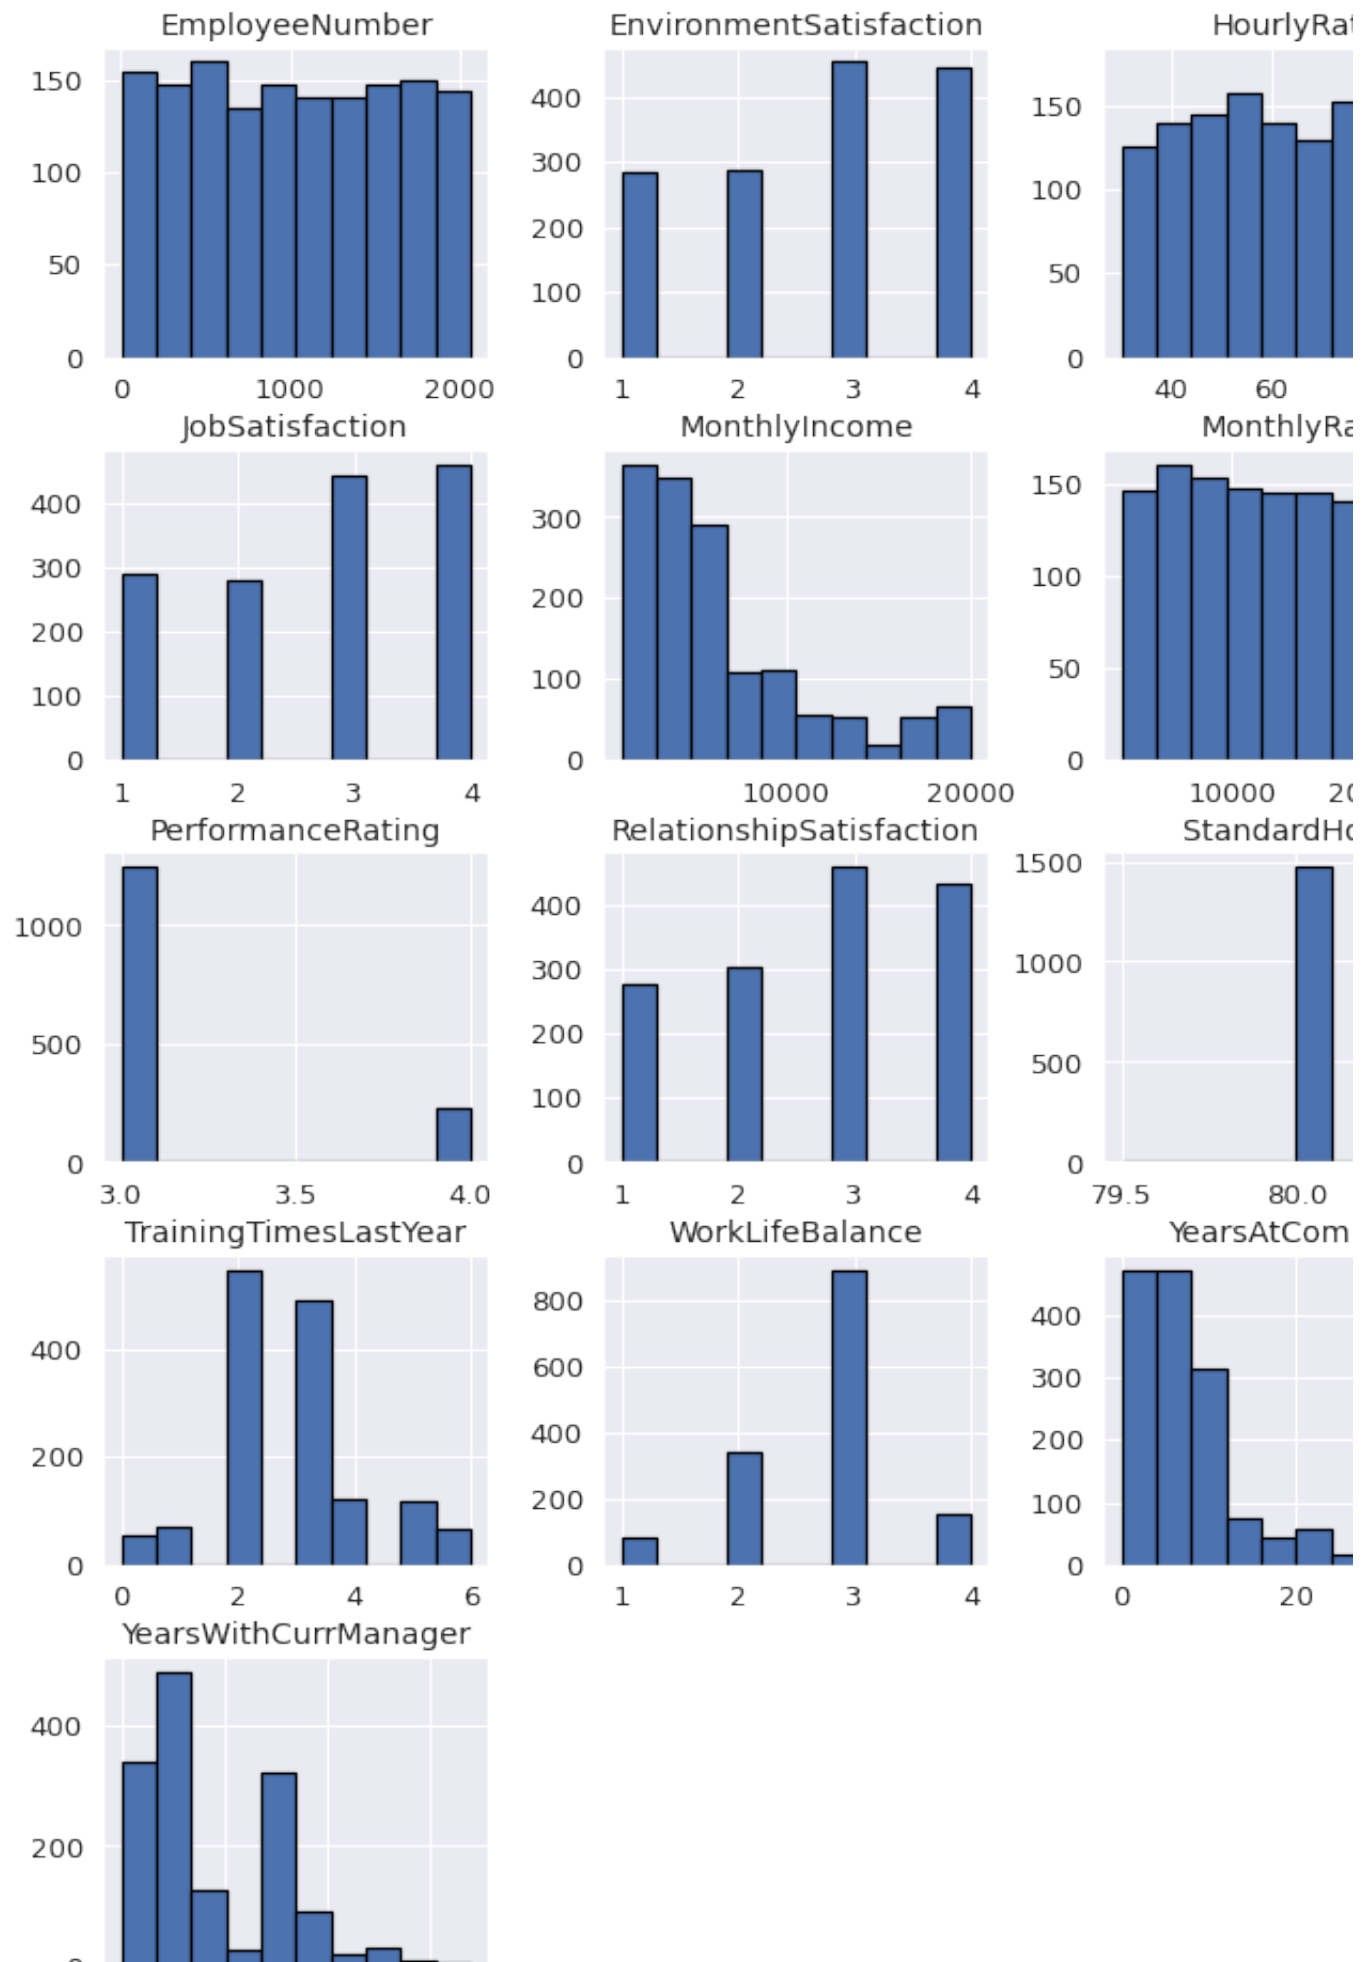
	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	En
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2

```
df.isnull().sum() #To check whether there are any missing values
```

```
Age 0
Attrition 0
BusinessTravel 0
DailyRate 0
Department 0
DistanceFromHome 0
Education 0
EducationField 0
EmployeeCount 0
EmployeeNumber 0
EnvironmentSatisfaction 0
Gender 0
HourlyRate 0
JobInvolvement 0
JobLevel 0
JobRole 0
JobSatisfaction 0
MaritalStatus 0
MonthlyIncome 0
MonthlyRate 0
NumCompaniesWorked 0
Over18 0
OverTime 0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
StandardHours 0
StockOptionLevel 0
TotalWorkingYears 0
TrainingTimesLastYear 0
WorkLifeBalance 0
YearsAtCompany 0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

```
df.hist(edgecolor='black', linewidth=1.2, figsize=(20, 20));
```







As we see that standard deviation for EmployeeCount, StandardHours, Over18 is near zero so we will drop it. Identification number which won't be helpful in prediction so we will also drop it.

```
df.drop(['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours'], axis="columns")
```

```
categorical_col = []
for column in df.columns:
    if df[column].dtype == object and len(df[column].unique()) <= 50:
        categorical_col.append(column)
        print(f"{column} : {df[column].unique()}")
        print("=====")
```

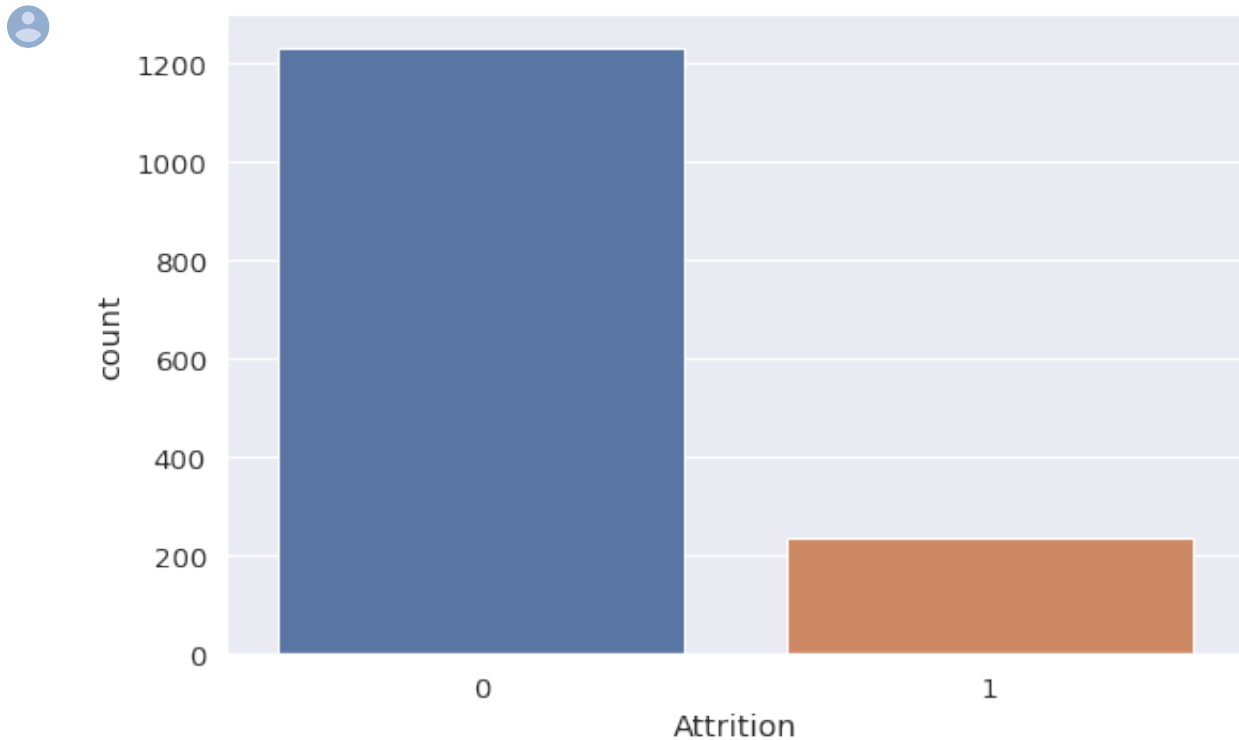
```
Attrition : ['Yes' 'No']
=====
BusinessTravel : ['Travel_Rarely' 'Travel_Frequently' 'Non-Travel']
=====
Department : ['Sales' 'Research & Development' 'Human Resources']
=====
EducationField : ['Life Sciences' 'Other' 'Medical' 'Marketing' 'Technical I
'Human Resources']
=====
Gender : ['Female' 'Male']
=====
JobRole : ['Sales Executive' 'Research Scientist' 'Laboratory Technician'
'Manufacturing Director' 'Healthcare Representative' 'Manager'
'Sales Representative' 'Research Director' 'Human Resources']
=====
MaritalStatus : ['Single' 'Married' 'Divorced']
=====
OverTime : ['Yes' 'No']
=====
```

```
df['Attrition'] = df.Attrition.astype("category").cat.codes
```

```
def plot_cat(attr, labels=None):
    if(attr=='JobRole'):
        sns.factorplot(data=df, kind='count', size=5, aspect=3, x=attr)
        return

    sns.factorplot(data=df, kind='count', size=5, aspect=1.5, x=attr)
```

```
plot_cat('Attrition')
```




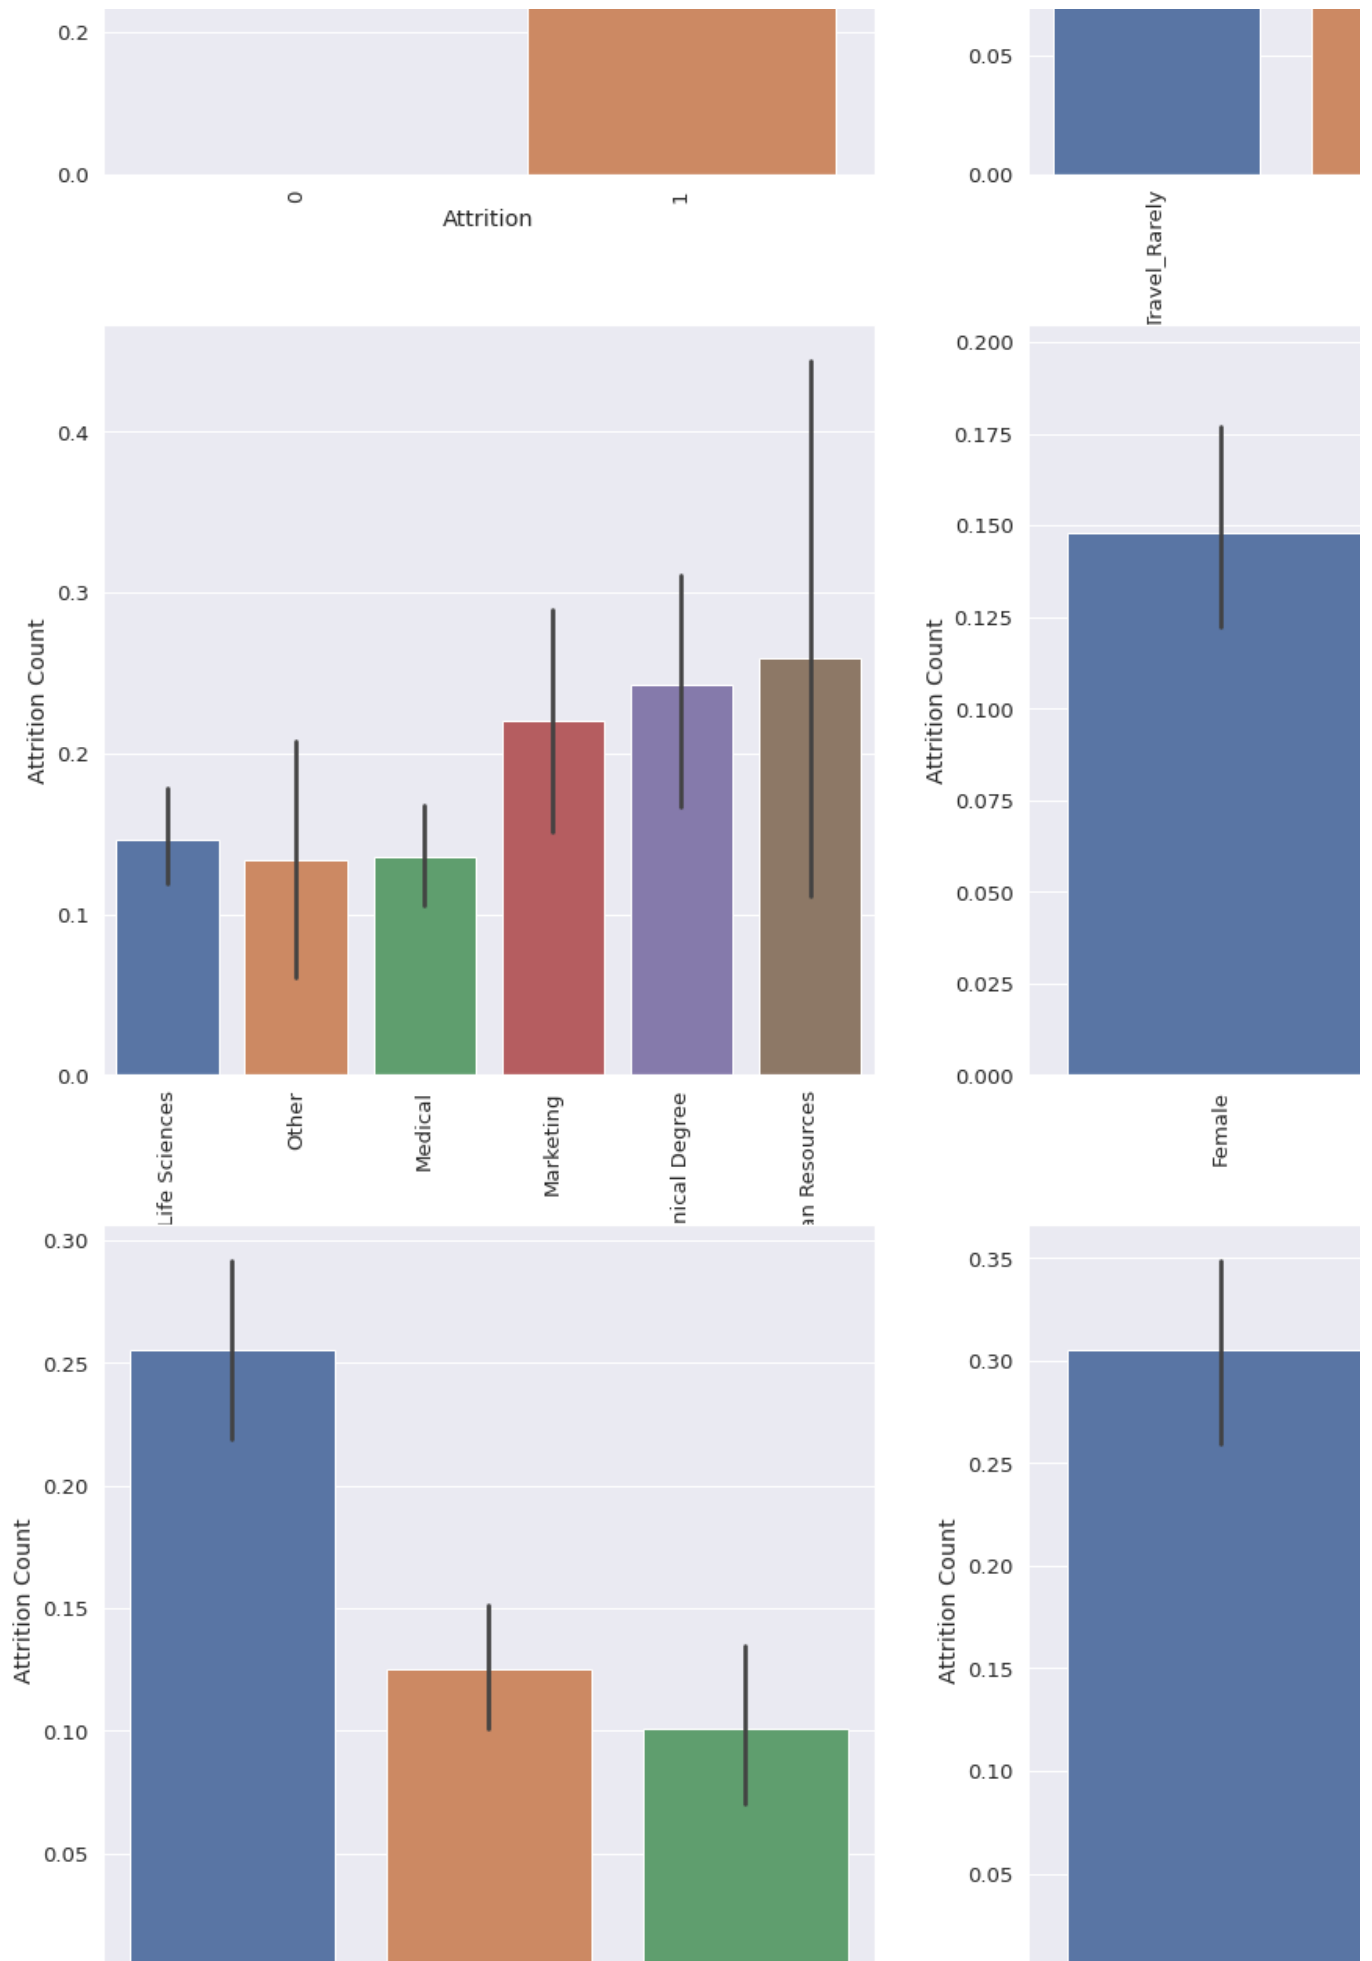
```
# Plotting how every feature correlate with the "target"
```


```
sns.set(font_scale=1.2)
```

```
plt.figure(figsize=(30, 30))
```

```
for i, column in enumerate(categorical_col, 1):  
    plt.subplot(3, 3, i)  
    g = sns.barplot(x=f"{column}", y='Attrition', data=df)  
    g.set_xticklabels(g.get_xticklabels(), rotation=90)  
    plt.ylabel('Attrition Count')  
    plt.xlabel(f'{column}')
```







▼ Conclusions

BusinessTravel : The workers who travel alot are more likely to quit then other employees.

Department : The worker in Research & Development are more likely to stay then the workers or

EducationField : The workers with Human Resources and Technical Degree are more likely to q
educations.

Gender : The Male are more likely to quit.

JobRole : The workers in Laboratory Technician, Sales Representative, and Human Resources a
positions.

MaritalStatus : The workers who have Single marital status are more likely to quit the Married, a

OverTime : The workers who work more hours are likely to quit then others.

```
plt.figure(figsize=(30, 30))
sns.heatmap(df.corr(), annot=True, cmap="RdYlGn", annot_kws={"size":15})
```



```
categorical_col.remove('Attrition')

# Transform categorical data into dummies
# categorical_col.remove("Attrition")
# data = pd.get_dummies(df, columns=categorical_col)
# data.info()
from sklearn.preprocessing import LabelEncoder

label = LabelEncoder()
for column in categorical_col:
    df[column] = label.fit_transform(df[column])
```



```

X = df.drop('Attrition', axis=1)
y = df.Attrition

from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, r

def print_score(clf, X, y, X_train, y_train, X_test, y_test, train=True):
    if train:
        pred = clf.predict(X_train)
        print("\nTrain Result:\n=====")
        print(f"accuracy score: {accuracy_score(y_train, pred):.4f}\n")
        train_accuracy.append(accuracy_score(y_train, pred))
        print(f"Classification Report: \n \tPrecision: {precision_score(y_train,
        print(f"Confusion Matrix: \n {confusion_matrix(y_train, clf.predict(X_tr

    elif train==False:
        pred = clf.predict(X_test)
        print("\nTest Result:\n=====")
        test_accuracy.append(accuracy_score(y_test, pred))
        print(f"accuracy score: {accuracy_score(y_test, pred)}\n")
        print(f"Classification Report: \n \tPrecision: {precision_score(y_test,
        print(f"Confusion Matrix: \n {confusion_matrix(y_test, pred)}\n")
        scores = cross_val_score(clf, X, y, cv=10)
        cv_accuracy.append(scores.mean())
        cv_devia.append(scores.std()*2)
        print(f"10-fold Cross Validation Accuracy: %0.2f (+/- %0.2f)" % (scores.

def plot_ROC(model, X_test, y_test):
    probs = model.predict_proba(X_test)
    preds = probs[:,1]
    fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
    roc_auc = metrics.auc(fpr, tpr)
    roc_accuracy.append(roc_auc)
    # method I: plt
    plt.title('Receiver Operating Characteristic')
    plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
    plt.legend(loc = 'lower right')
    plt.plot([0, 1], [0, 1], 'r--')
    plt.xlim([0, 1])
    plt.ylim([0, 1])
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()

```

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_

train_accuracy = []
test_accuracy = []
cv_accuracy = []
cv_devia = []
roc_accuracy = []
```

► Decision Tree

↳ 5 cells hidden

► Random Forest

↳ 7 cells hidden

► Logistic Regression

↳ 3 cells hidden

► SVM

↳ 7 cells hidden

