

Problem statement 1

```
import matplotlib.pyplot as plt
from matplotlib import style
import numpy as np
#from sklearn import preprocessing, cross_validation
import pandas as pd
```

```
import numpy as np
```

```
np.random.seed(123)
```


```
allwalks = []
```

```
for i in range(250):
    randwalk = [0]
    for x in range(100):
        step = randwalk[-1]
        dice = np.random.randint(1,7)
        if dice <= 2 :
            step = max(0, step - 1)

        elif dice<=5:
            step += 1

        else:
            step = step + np.random.randint(1,7)

    print(step)
```



```
1
2
0
0
0
0
1
1
0
1
1
1
0
0
5
1
^
```

[illegible]

1
4
0
0
0
0
0
1
1
1
0
1
1
1
1
4
0
1
1
1
1
0
0
0
0
0
1
1
1
0
2
1
0
6
0
0
1
6
0
3
4
1
1
6
1
0
5
1
0
1
5
0
1
1
1
0
~

0
1
1
5
1
1
1
1
0
0
0
0
5
0
0
0
0
1
4
1
1
1
1
1
1
1
0
1
1
1
0
1
1
1
1
1
0
1
1
1
1
0
1
1
3
3
1
1
0
1
0
1
1
1
1
1
1
1
1
1
1
0
0
1
-

0
0
0
1
0
0
6
1
2
0
1
1
1
1
3
1
1
1
1
1
1
0
1
1
1
1
1
1
0
1
0
4
1
1
1
1
1
1
0
1
0
0
1
0
1
1
1
1
1
1
1
1
1
0
0
0
1
1
1
1
1
1
1
1
1
0
0
1
2
0
0
0

```

0
1
1
0
4
1
1
1
4
1
2
1
0
1
1
0
1
1

```

Problem statement 2

Random data for multiple linear regression

```

import numpy as np
import pandas as pd
import scipy
import random
from scipy.stats import norm
random.seed(1)
n_features = 4
X = []
for i in range(n_features):
    X_i = scipy.stats.norm.rvs(0, 1, 100)
    X.append(X_i)
#print(X)
eps = scipy.stats.norm.rvs(0, 0.25, 100)
y = 1 + (0.4 * X[0]) + eps + (0.5 * X[1]) + (0.3 * X[2]) + (0.4 * X[3])
data_mlr = {'X0': X[0], 'X1':X[1], 'X2':X[2], 'X3':X[3], 'Y': y }
df = pd.DataFrame(data_mlr)
print(df.head())
print(df.tail())
print(df.info())
print(df.describe())
#df.to_csv('file1.csv')

```



```

      X0      X1      X2      X3      Y
0  1.229187 -0.272740 -2.976897  0.234832  0.514580
1 -0.504641  0.137005  1.243000  0.360801  1.094348
2 -1.716661 -0.621364 -0.874199 -0.667714 -0.531753
3 -0.401698 -0.171870  0.105726  0.719925  1.319178
4  2.338396 -0.389372  0.371695 -0.211898  2.015665
      X0      X1      X2      X3      Y
95  0.538347  0.947229  1.702177  1.827932  2.946796
96 -2.124628 -0.023599 -0.737292  0.959677  0.382169
97  0.690109 -1.404677  1.814179 -1.350687  0.695210
98 -0.473165  0.264512 -0.011203  2.202028  1.598639
99  0.272442  0.936071 -1.283808  1.368407  1.735643
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  ------  -
0    X0      100 non-null    float64
1    X1      100 non-null    float64
2    X2      100 non-null    float64
3    X3      100 non-null    float64
4    Y        100 non-null    float64
dtypes: float64(5)
memory usage: 4.0 KB
None

```

	X0	X1	X2	X3	Y
count	100.000000	100.000000	100.000000	100.000000	100.000000
mean	-0.152297	0.030112	-0.063573	0.006363	0.978705
std	0.924978	0.975453	1.096232	1.060847	0.913494
min	-2.206992	-2.118134	-2.976897	-2.946925	-1.051042
25%	-0.729312	-0.621701	-0.698830	-0.674253	0.377647
50%	-0.158125	0.089349	-0.141087	-0.037818	0.964189
75%	0.498398	0.641858	0.574192	0.620150	1.549381
max	2.338396	3.159726	3.210134	2.292966	3.643752

Random data for logistic regression

```

n_features = 4
X = []
for i in range(n_features):
    X_i = scipy.stats.norm.rvs(0, 1, 100)
    X.append(X_i)
#print(X)
a1 = (np.exp(1 + (0.5 * X[0]) + (0.4 * X[1]) + (0.3 * X[2]) + (0.5 * X[3])))/(1 +
#print(a1)
y1 = []
for i in a1:
    if (i>=0.5):
        y1.append(1)
    else:

```

```

y1.append(0)
#print(y1)
data_lr = {'X0': X[0], 'X1':X[1], 'X2':X[2], 'X3':X[3], 'Y': y1 }
df1 = pd.DataFrame(data_lr)
print(df.head())
print(df.tail())
print(df.info())
print(df.describe())

```



```

      X0      X1      X2      X3      Y
0  0.310326 -0.538983  0.522009 -0.630752  0.888074
1  0.026933  1.005510 -1.519784  0.317596  1.706252
2  1.454472 -1.948507  0.989502  1.673824  2.006770
3  0.299680 -1.090324 -0.968199  0.285824  0.376355
4  1.568637  0.042656 -0.204593  1.126121  2.629879
      X0      X1      X2      X3      Y
95 -0.408410  0.615359 -2.553963  1.017630  0.665036
96  1.271942  0.739803  1.451475 -2.180999  1.219121
97 -1.462029 -1.407781  0.657375  0.320367  0.309101
98 -0.355275 -1.795455  0.762725 -0.701842 -0.118037
99 -0.845194 -0.817530 -1.009229  0.328676 -0.005031
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0    X0      100 non-null    float64
1    X1      100 non-null    float64
2    X2      100 non-null    float64
3    X3      100 non-null    float64
4    Y       100 non-null    float64
dtypes: float64(5)
memory usage: 4.0 KB
None

```

	X0	X1	X2	X3	Y
count	100.000000	100.000000	100.000000	100.000000	100.000000
mean	0.012668	-0.098384	0.002392	0.037210	1.007865
std	0.984979	1.056384	0.890788	1.027167	0.875303
min	-2.174584	-2.241255	-2.553963	-3.162631	-0.881631
25%	-0.761235	-0.900435	-0.565406	-0.652427	0.344661
50%	-0.013255	-0.008498	-0.023229	0.078830	1.035789
75%	0.724811	0.740714	0.684799	0.733959	1.709912
max	2.259437	2.192720	2.373255	2.320635	2.909347

Random data for K means clustering


```

X_a= -2 * np.random.rand(100,2)
X_b = 1 + 2 * np.random.rand(50,2)
X_a[50:100, :] = X_b
plt.scatter(X_a[:, 0], X_a[:, 1], s = 50)

```



```
plt.show()
data_kmeans = {'X0': X_a[:,0], 'X1':X_a[:,1]}
df3 = pd.DataFrame(data_kmeans)
print(df.head())
print(df.tail())
print(df.info())
print(df.describe())
```



```

      X0      X1      X2      X3      Y
0  0.310326 -0.538983  0.522009 -0.630752  0.888074
1  0.026933  1.005510 -1.519784  0.317596  1.706252
2  1.454472 -1.948507  0.989502  1.673824  2.006770
3  0.299680 -1.090324 -0.968199  0.285824  0.376355
4  1.568637  0.042656 -0.204593  1.126121  2.629879

```

```

      X0      X1      X2      X3      Y
95 -0.408410  0.615359 -2.553963  1.017630  0.665036
96  1.271942  0.739803  1.451475 -2.180999  1.219121
97 -1.462029 -1.407781  0.657375  0.320367  0.309101
98 -0.355275 -1.795455  0.762725 -0.701842 -0.118037
99 -0.845194 -0.817530 -1.009229  0.328676 -0.005031

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100 entries, 0 to 99
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	X0	100 non-null	float64
1	X1	100 non-null	float64
2	X2	100 non-null	float64
3	X3	100 non-null	float64
4	Y	100 non-null	float64

```
dtypes: float64(5)
```

```
memory usage: 4.0 KB
```

```
None
```

	X0	X1	X2	X3	Y
count	100.000000	100.000000	100.000000	100.000000	100.000000
mean	0.012668	-0.098384	0.002392	0.037210	1.007865
std	0.984979	1.056384	0.890788	1.027167	0.875303
min	-2.174584	-2.241255	-2.553963	-3.162631	-0.881631
25%	-0.761235	-0.900435	-0.565406	-0.652427	0.344661
50%	-0.013255	-0.008498	-0.023229	0.078830	1.035789
75%	0.724811	0.740714	0.684799	0.733959	1.709912
max	2.259437	2.192720	2.373255	2.320635	2.909347

Problem statement - 3

Linear Regression using gradient descent

```
X = df.iloc[:,0].values
#print(X)
y = df.iloc[:,4].values
b1 = 0
b0 = 0
l = 0.001
epochs = 100

n = float(len(X))
for i in range(epochs):
    y_p = b1*X + b0
    loss = np.sum(y_p - y1)**2
    d1 = (-2/n) * sum(X * (y - y_p))
    d0 = (-2/n) * sum(y - y_p)
    b1 = b1 - (l*d1)
    b0 = b0 - (l*d0)

print(b1,b0)
```



Logistic regression using Gradient descent

```
X1 = df1.iloc[:,0:4].values
y1 = df1.iloc[:,4].values

def sigmoid(Z):
    return 1 / (1 + np.exp(-Z))

def loss(y1, y_hat):
    return -np.mean(y1 * np.log(y_hat) + (1 - y1) * (np.log(1 - y_hat)))

W = np.zeros((4, 1))
b = np.zeros((1, 1))

m = len(y1)
lr = 0.001
for epoch in range(1000):
    Z = np.matmul(X1, W) + b
    A = sigmoid(Z)
    logistic_loss = loss(y1, A)
    dz = A - y1
    dw = 1/m * np.matmul(X1.T, dz)
    db = np.sum(dz)

    W = W - lr * dw
    b = b - lr * db

    if epoch % 100 == 0:
        print(logistic_loss)
```



Linear Regreesion using L1 Regularization

```
X = df.iloc[:,0].values
#print(X)
y = df.iloc[:,4].values
b1 = 0
b0 = 0
l = 0.001
epochs = 100
lam = 0.1

n = float(len(X))
for i in range(epochs):
    y_p = b1*X + b0
    loss = np.sum(y_p - y)**2 + (lam * b1)
    d1 = (-2/n) * sum(X * (y - y_p)) + lam
    d0 = (-2/n) * sum(y - y_p)
    b1 = b1 - (l*d1)
    b0 = b0 - (l*d0)

print(b1,b0)
```



Linear Regreesion using L2 Regularization

```
X = df.iloc[:,0].values
#print(X)
y = df.iloc[:,4].values
b1 = 0
b0 = 0
l = 0.001
epochs = 100
lam = 0.1

n = float(len(X))
for i in range(epochs):
    y_p = b1*X + b0
    loss = np.sum(y_p - y)**2 + ((lam/2) * b1)
    d1 = (-2/n) * sum(X * (y - y_p)) + (lam * b1)
    d0 = (-2/n) * sum(y - y_p)
    b1 = b1 - (l*d1)
    b0 = b0 - (l*d0)

print(b1,b0)
```



Logistic regression using L1 regularization

```
X1 = df1.iloc[:,0:4].values
y1 = df1.iloc[:,4].values
lam = 0.1
def sigmoid(Z):
    return 1 / (1 + np.exp(-Z))

def loss(y1, y_hat):
    return -np.mean(y1 * np.log(y_hat) + (1 - y1) * (np.log(1 - y_hat))) + (lam * (np.sum(

W = np.zeros((4, 1))
b = np.zeros((1, 1))

m = len(y1)
lr = 0.001
for epoch in range(1000):
    Z = np.matmul(X1, W) + b
    A = sigmoid(Z)
    logistic_loss = loss(y1, A)
    dz = A - y1
    dw = 1/m * np.matmul(X1.T, dz) + lam
    db = np.sum(dz)

    W = W - lr * dw
    b = b - lr * db

    if epoch % 100 == 0:
        print(logistic_loss)
```



Logistic regression using L2 regularization

```
X1 = df1.iloc[:,0:4].values
y1 = df1.iloc[:,4].values
lam = 0.1
def sigmoid(Z):
    return 1 / (1 + np.exp(-Z))

def loss(y1, y_hat):
    return -np.mean(y1 * np.log(y_hat) + (1 - y1) * (np.log(1 - y_hat))) + (lam * (np.sum(

W = np.zeros((4, 1))
b = np.zeros((1, 1))

m = len(y1)
lr = 0.001
for epoch in range(1000):
    Z = np.matmul(X1, W) + b
    A = sigmoid(Z)
    logistic_loss = loss(y1, A)
    dz = A - y1
    dw = 1/m * np.matmul(X1.T, dz) + lam * W
    db = np.sum(dz)

    W = W - lr * dw
    b = b - lr * db

    if epoch % 100 == 0:
        print(logistic_loss)
```



K Means Clustering Algorithm

```
class K_Means:
    def __init__(self, k=2, tol=0.001, max_iter=300):
```



```

self.k = k
self.tol = tol
self.max_iter = max_iter

def fit(self, data):

    self.centroids = {}

    for i in range(self.k):
        self.centroids[i] = data[i]

    for i in range(self.max_iter):
        self.classifications = {}

        for i in range(self.k):
            self.classifications[i] = []

        for featureset in X:
            distances = [np.linalg.norm(featureset-self.centroids[centroid]) for centroid in self.centroids]
            classification = distances.index(min(distances))
            self.classifications[classification].append(featureset)

        prev_centroids = dict(self.centroids)

        for classification in self.classifications:
            self.centroids[classification] = np.average(self.classifications[classification], axis=0)

        optimized = True

        for c in self.centroids:
            original_centroid = prev_centroids[c]
            current_centroid = self.centroids[c]
            if np.sum((current_centroid-original_centroid)/original_centroid) > self.tol:
                print(np.sum((current_centroid-original_centroid)/original_centroid))
                optimized = False

        if optimized:
            break

def predict(self, data):
    distances = [np.linalg.norm(data-self.centroids[centroid]) for centroid in self.centroids]
    classification = distances.index(min(distances))
    return classification

colors = 10*["g", "r", "c", "b", "k"]

```

```
X = df3.iloc[:,0:2].values
clf = K_Means()
clf.fit(X)

for centroid in clf.centroids:
    plt.scatter(clf.centroids[centroid][0], clf.centroids[centroid][1],
                marker="o", color="k", s=150, linewidths=5)

for classification in clf.classifications:
    color = colors[classification]
    for featureset in clf.classifications[classification]:
        plt.scatter(featureset[0], featureset[1], marker="x", color=color, s=150)
```



Problem Statement - 4

Linear Regression from scratch using OOPS

```
import numpy as np

class LinearRegressionModel():

    def __init__(self, dataset, learning_rate, num_iterations):
        self.dataset = np.array(dataset)
        self.b = 0
        self.m = 0
        self.learning_rate = learning_rate
        self.num_iterations = num_iterations
        self.M = len(self.dataset)
        self.total_error = 0
```

```

def apply_gradient_descent(self):
    for i in range(self.num_iterations):
        self.do_gradient_step()

def do_gradient_step(self):
    b_summation = 0
    m_summation = 0
    for i in range(self.M):
        x_value = self.dataset[i, 0]
        y_value = self.dataset[i, 1]
        b_summation += (((self.m * x_value) + self.b) - y_value)
        m_summation += (((self.m * x_value) + self.b) - y_value) * x_value
    self.b = self.b - (self.learning_rate * (1/self.M) * b_summation)
    self.m = self.m - (self.learning_rate * (1/self.M) * m_summation)

def compute_error(self):
    for i in range(self.M):
        x_value = self.dataset[i, 0]
        y_value = self.dataset[i, 1]
        self.total_error += ((self.m * x_value) + self.b) - y_value
    return self.total_error

def __str__(self):
    return "Results: b: {}, m: {}, Final Total error: {}".format(round(self.

def get_prediction_based_on(self, x):
    return round(float((self.m * x) + self.b), 2) # Type: Numpy float.

def main():
    school_dataset = np.genfromtxt(DATASET_PATH, delimiter=",")
    lr = LinearRegressionModel(school_dataset, 0.0001, 1000)
    lr.apply_gradient_descent()
    hours = [10, 20, 30, 40, 50, 60, 70, 80, 90, 100]
    for hour in hours:
        print("Studied {} hours and got {} points.".format(hour, lr.get_predicti
    print(lr)

if __name__ == "__main__": main()

```



* *italicized text* Logistic Regression from scratch using OOPS**

```

class LogisticRegression:
    def __init__(self, learning_rate, num_iters, fit_intercept = True, verbose = F
        self.learning_rate = learning_rate
        self.num_iters = num_iters
        self.fit_intercept = fit_intercept
        self.verbose = verbose
    def __add_intercept(self, X):
        intercept = np.ones((X.shape[0],1))
        return np.concatenate((intercept,X),axis=1)
    def __sigmoid(self,z):
        return 1/(1+np.exp(-z))
    def __loss(self, h, y):
        return (-y * np.log(h) - (1-y) * np.log(1-h)).mean()

    def fit(self,X,y):
        if self.fit_intercept:
            X = self.__add_intercept(X)
            self.theta = np.zeros(X.shape[1])

        for i in range(self.num_iters):
            z = np.dot(X,self.theta)
            h = self.__sigmoid(z)
            gradient = np.dot(X.T,(h-y))/y.size

            self.theta -= self.learning_rate * gradient

            z = np.dot(X,self.theta)
            h = self.__sigmoid(z)
            loss = self.__loss(h,y)

            if self.verbose == True and i % 1000 == 0:
                print(f'Loss: {loss}\t')
    def predict_probability(self,X):
        if self.fit_intercept:
            X = self.__add_intercept(X)
            return self.__sigmoid(np.dot(X,self.theta))
    def predict(self,X):
        return (self.predict_probability(X).round())

```

K Means from scratch using OOPS

```

class K_Means:
    def __init__(self, k=2, tol=0.001, max_iter=300):
        self.k = k
        self.tol = tol
        self.max_iter = max_iter

```

```

def fit(self,data):

    self.centroids = {}

    for i in range(self.k):
        self.centroids[i] = data[i]

    for i in range(self.max_iter):
        self.classifications = {}

        for i in range(self.k):
            self.classifications[i] = []

        for featureset in X:
            distances = [np.linalg.norm(featureset-self.centroids[centroid]) for centroid in self.centroids]
            classification = distances.index(min(distances))
            self.classifications[classification].append(featureset)

        prev_centroids = dict(self.centroids)

        for classification in self.classifications:
            self.centroids[classification] = np.average(self.classifications[classification], axis=0)

        optimized = True

        for c in self.centroids:
            original_centroid = prev_centroids[c]
            current_centroid = self.centroids[c]
            if np.sum((current_centroid-original_centroid)/original_centroid) != 0:
                print(np.sum((current_centroid-original_centroid)/original_centroid))
                optimized = False

        if optimized:
            break

    def predict(self,data):
        distances = [np.linalg.norm(data-self.centroids[centroid]) for centroid in self.centroids]
        classification = distances.index(min(distances))
        return classification

colors = 10*["g","r","c","b","k"]

```

